



北京大学心理与认知科学学院

School of Psychological and Cognitive Sciences, Peking University

Computational modeling in RStan --basics and practices

Hanbo Xie



- ❑ Name: Hanbo Xie
- ❑ Current Affiliation: Research Assistant in Zhou's Lab in Peking University
- ❑ Education: BSc of Human Resource Management in SWUFE; Exchange student in UCR
- ❑ Research Interests: Interested in decision making, learning (reinforcement learning, social learning and group learning) via computational neuroscience approach.





- This tutorial is not to:
 - Introduce how to create a good model;
 - Introduce a specific family of models;
- But we expect to:
 - Develop an interest in combining computational models in your own research
 - Know a basic routine to conduct computational model work
 - Use MCMC and Rstan to fit your own models



Basics about computational modeling

Standard approaches of modeling

MCMC and Rstan


Practices: implementing a reinforcement learning
model in Rstan



Basics about computational modeling



- What is computational modeling?
 - computational modeling is a technique that applies **mathematical models** and **computer science** in complex systems (e.g., ecological environment, biology, physics and etc.,) to better understand the underlying mechanisms.
- What can computational modeling relate to psychology and neuroscience?
 - better understand behavioral mechanisms (choice behavior, RT, eye movement and etc.)
 - better understand neural mechanisms (DCM, neuronal spikes and etc.)
- Why is it special?
 - high precision on predicting data
 - inference on latent variables (states)
 - capability of coping large-scale, high-dimensional data
 - ...



Computational cognitive research is growing rapidly, due to the development of machine learning techniques and AI.

New interdisciplinary fields emerged:

- neuroeconomics
- computational neuroscience
- computational psychiatry
- ...



- Basic concepts of a computational model:
 - suppose you have a set of data concerning RT and choice accuracy....

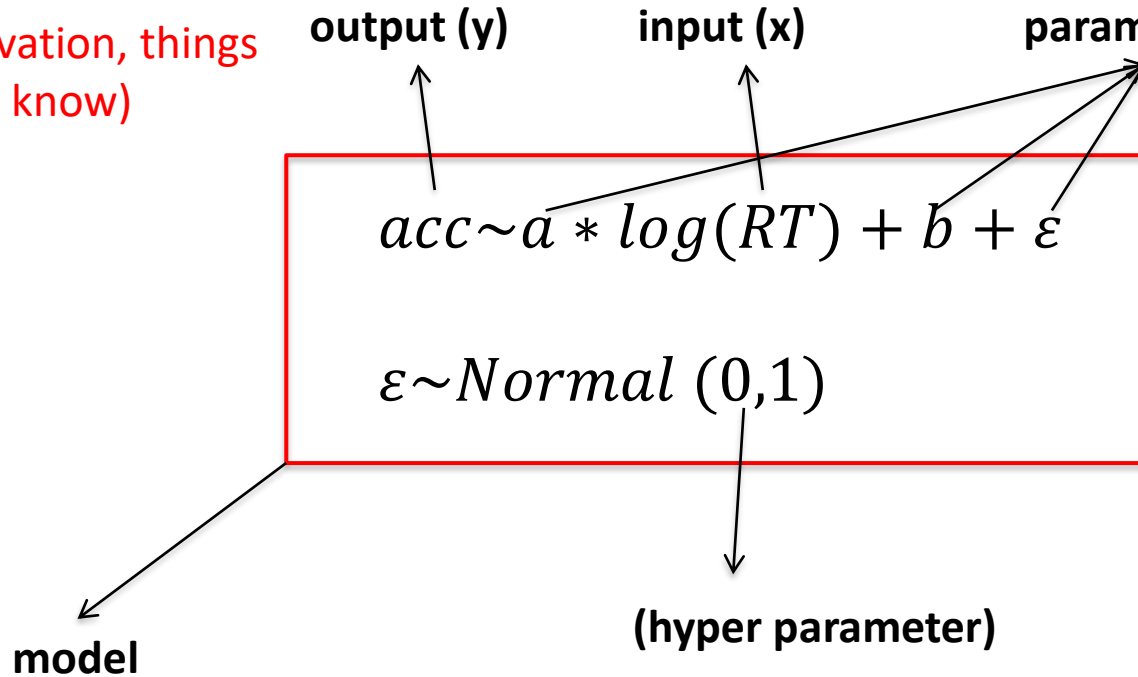
data (observation, things
you already know)

output (y)

input (x)

parameters

goal (things
you wish to
know...)



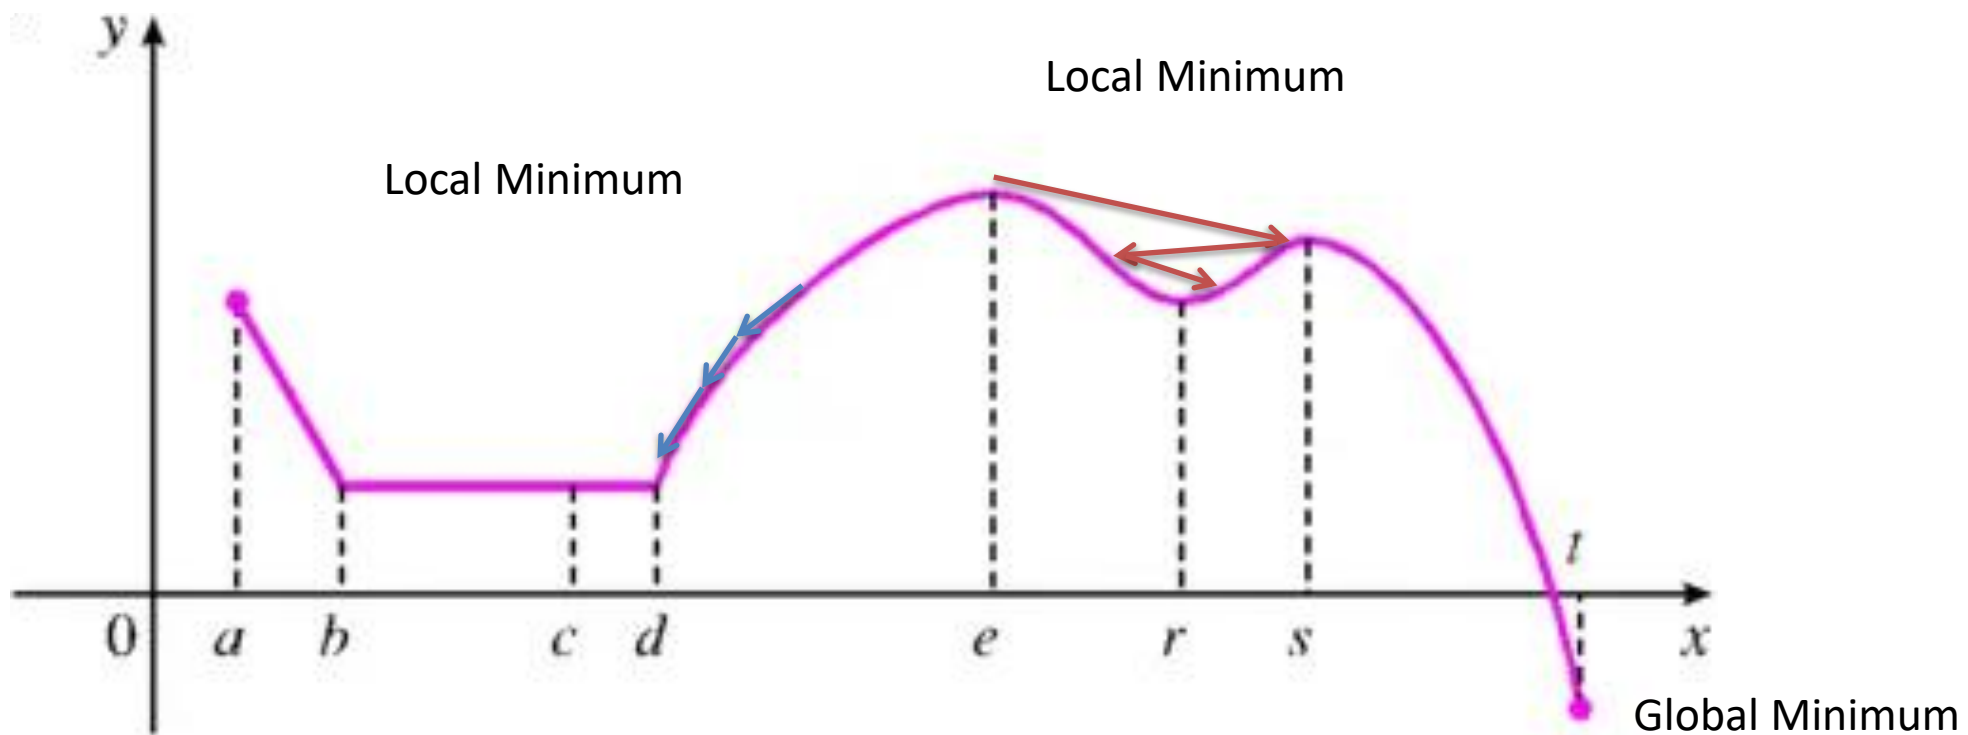
relationship (hypothesis,
things you think you know)




- Model fitting:
 - get to the goal
 - optimize the parameters that best suit to your data, given your model
 - how to optimize: minimize the errors between data and model predictions, by adjusting the parameters.
- Cost functions (measurement of ‘errors’):
 - mean squared error (MSE)
 - log-likelihood
 - TD error (in deep reinforcement learning)
 - dis-similarities
 - ...




- Two general approaches of estimating the parameters:
 - iterative approaches
 - » Gradient descent (MLE), Variational Bayesian Analysis, (mean-field, newton, ...)
 - » Quick computing, but may reach a local minimum
 - sampling approaches
 - » MCMC, Gibson sampling, Grid search...
 - » Computationally expensive, hard to sample for extreme complex models







残差





First order polynomial

(Linear regression)

$$y = \theta_1 x + \theta_0$$

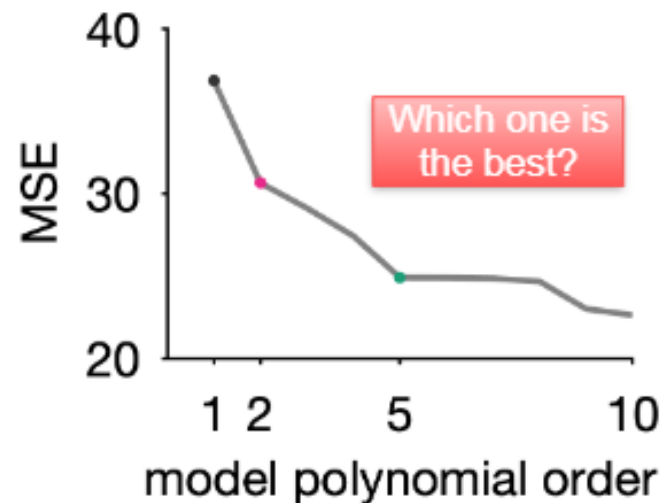
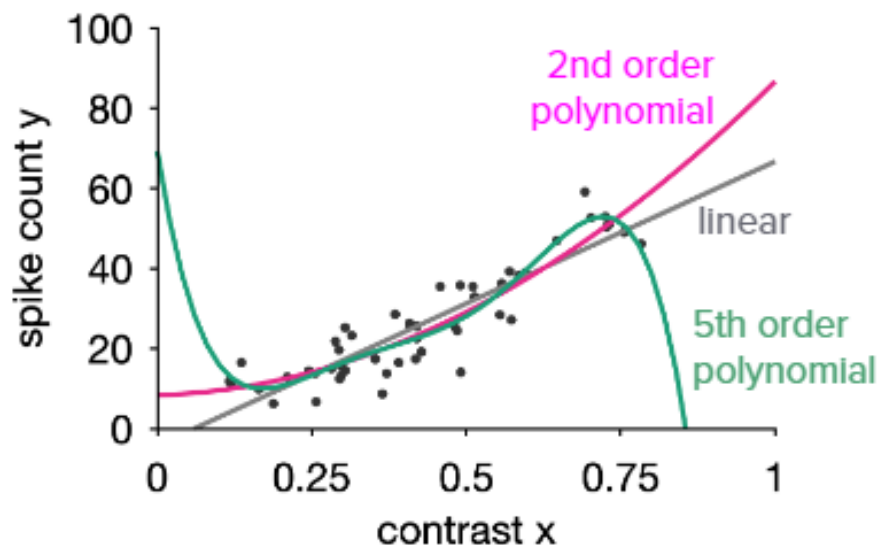
Second order polynomial

$$y = \theta_2 x^2 + \theta_1 x + \theta_0$$

Higher order



polynomial

$$y = \sum_{p=0}^5 \theta_p x^p$$



Overfitting!







- How to avoid overfitting?
 - Train with more data!
 - Machine learning:
 - Regularization
 - Cross-validation
 - Dropout
 - ...
 - Psychology:
 - Model comparison (a simpler model)
 - Information criteria (AIC, BIC, DIC, LOOIC...)



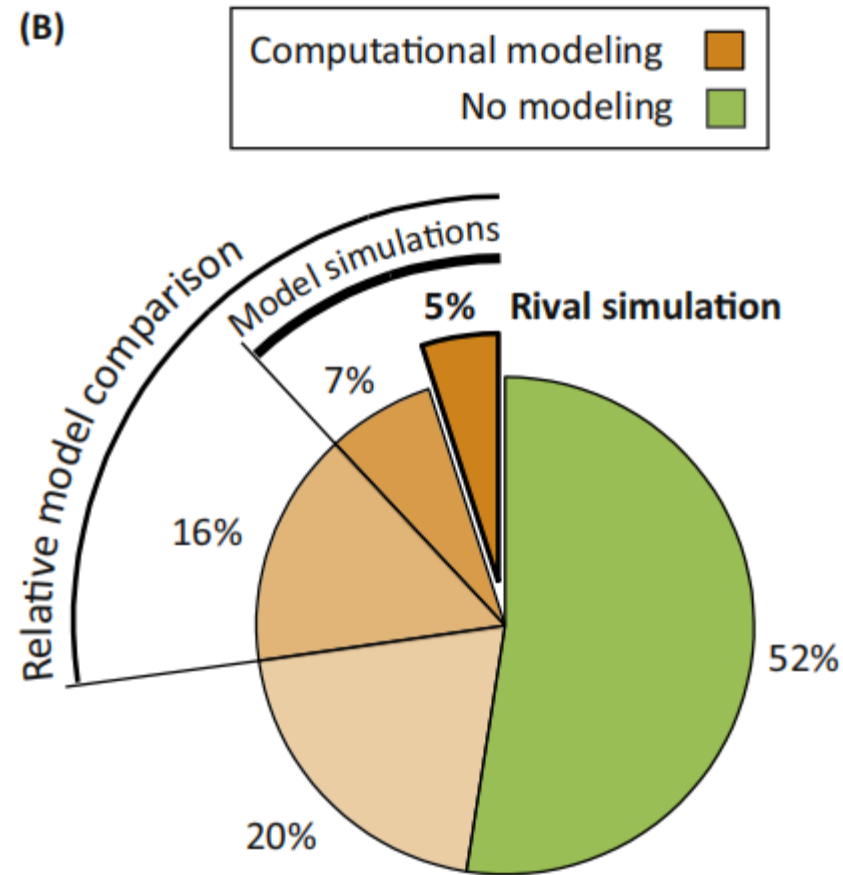
- What is computational modeling? Why do we need computational modeling?
- Basic concepts of a model
- What is model fitting? Two general approaches for fitting.
- What is overfitting? Why do we need to avoid overfitting?



Standard approaches of modeling



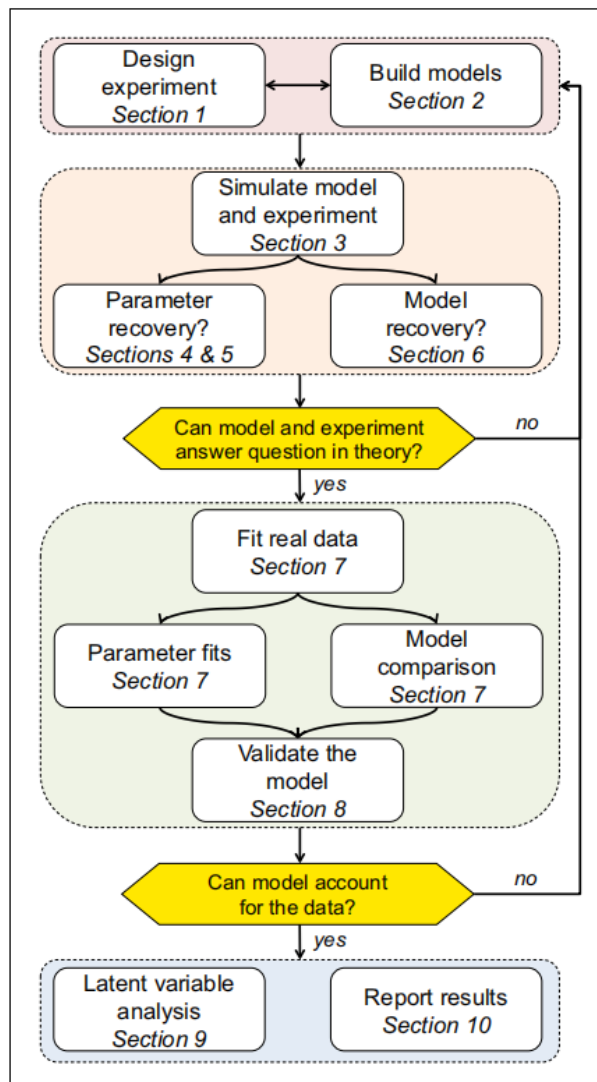
- Like all other tools, computational modeling requires golden standards to be used.



Palmini, et.al. 2017



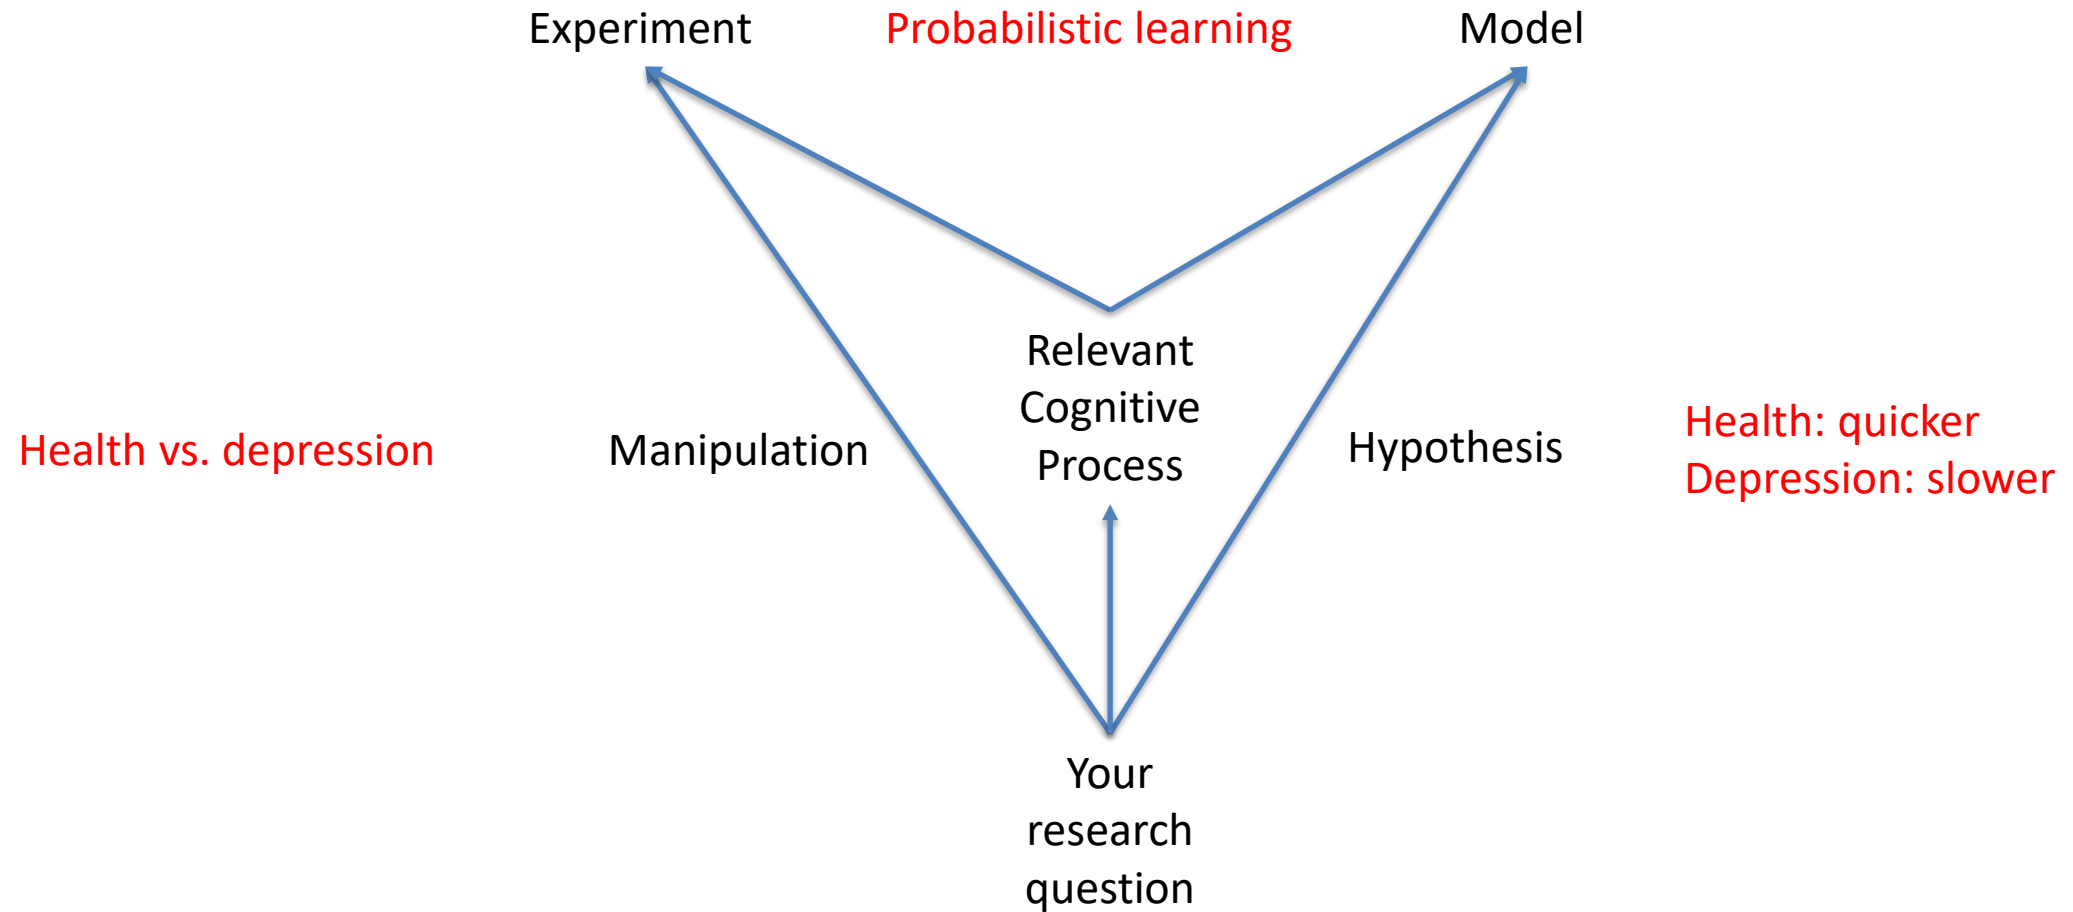
Ten Simple Rules



Wilson, Colins, 2019



Step 1: experiment vs. model



e.g. What is difference between healthy people and those with depression in probabilistic learning?



Model \longrightarrow Fake Data

Pre-set parameter

Health (N=30)

Depression (N=30)

Model: Reinforcement learning model

Environment(Task): Probability reversal learning task

$$PE = R - v_{i,t}$$

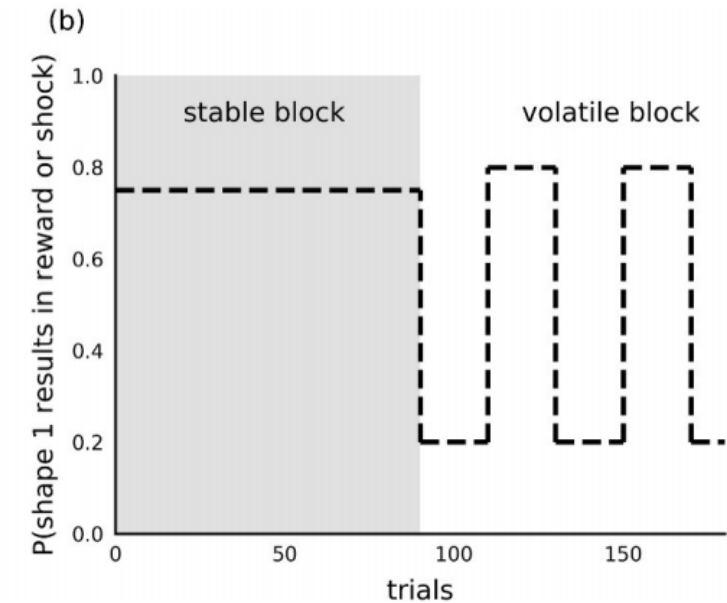
Prediction Error

$$v_{i,t+1} = v_{i,t} + \alpha * PE$$

α : learning rate

$$P(i, t + 1) = \frac{e^{v_{i,t+1}}}{\sum e^{v_{i,t+1}}}$$

Rescorla, R. A., & Wagner, A. R. (1972)





Step-by-Step Simulation

1. Define the environment

```
#adopt WSLS strategy  
p_reward1=c(0.8,0.2)  
p_reward2=c(0.2,0.8)
```

2. Define (an) agent(s)

- Action(choice)
- Reward(feedback)
- Policy (strategy)
- Hidden States...

```
if (t==1){  
  choice[t]=sample(c(1:2), size = 1, replace = T, prob = p_choose)
```

```
v[choice[t]]=v[choice[t]]+lr*(tmp_reward-v[choice[t]])
```

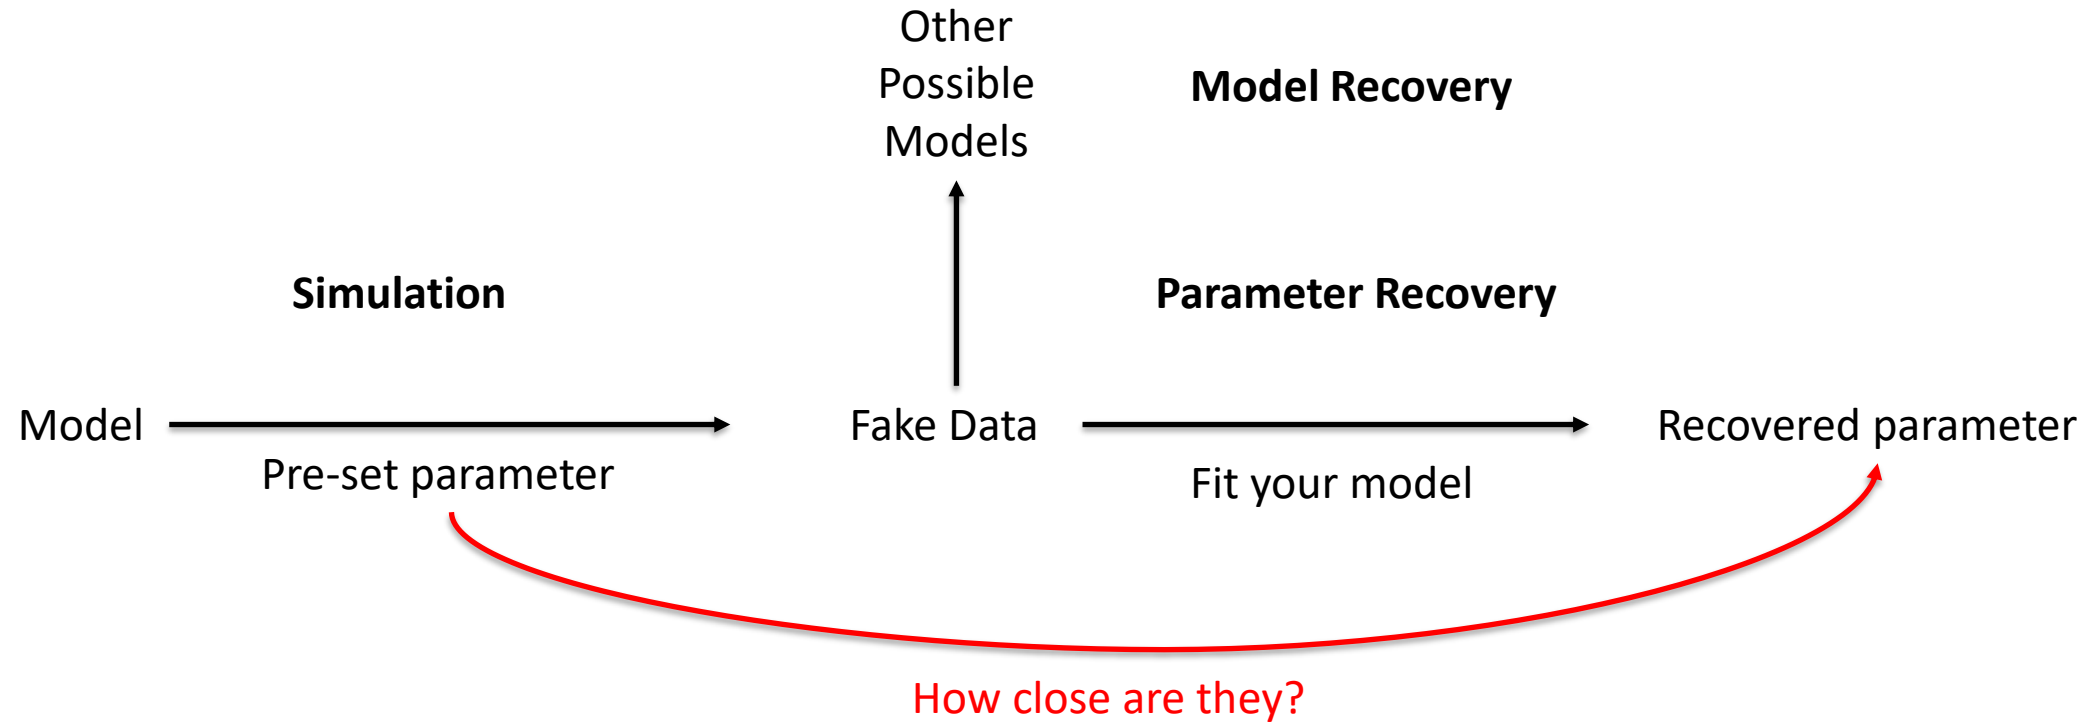
3. Define the environment response

```
tmp_reward=sample(c(1,-1), size = 1, replace = T, prob = p_reward[choice[t],])  
accumulated_reward2[t]=accumulated_reward2[t-1]+tmp_reward
```

4. Set and run the simulation process!

for-loop: trials and agents

5. Analyze and visualize your simulated data!

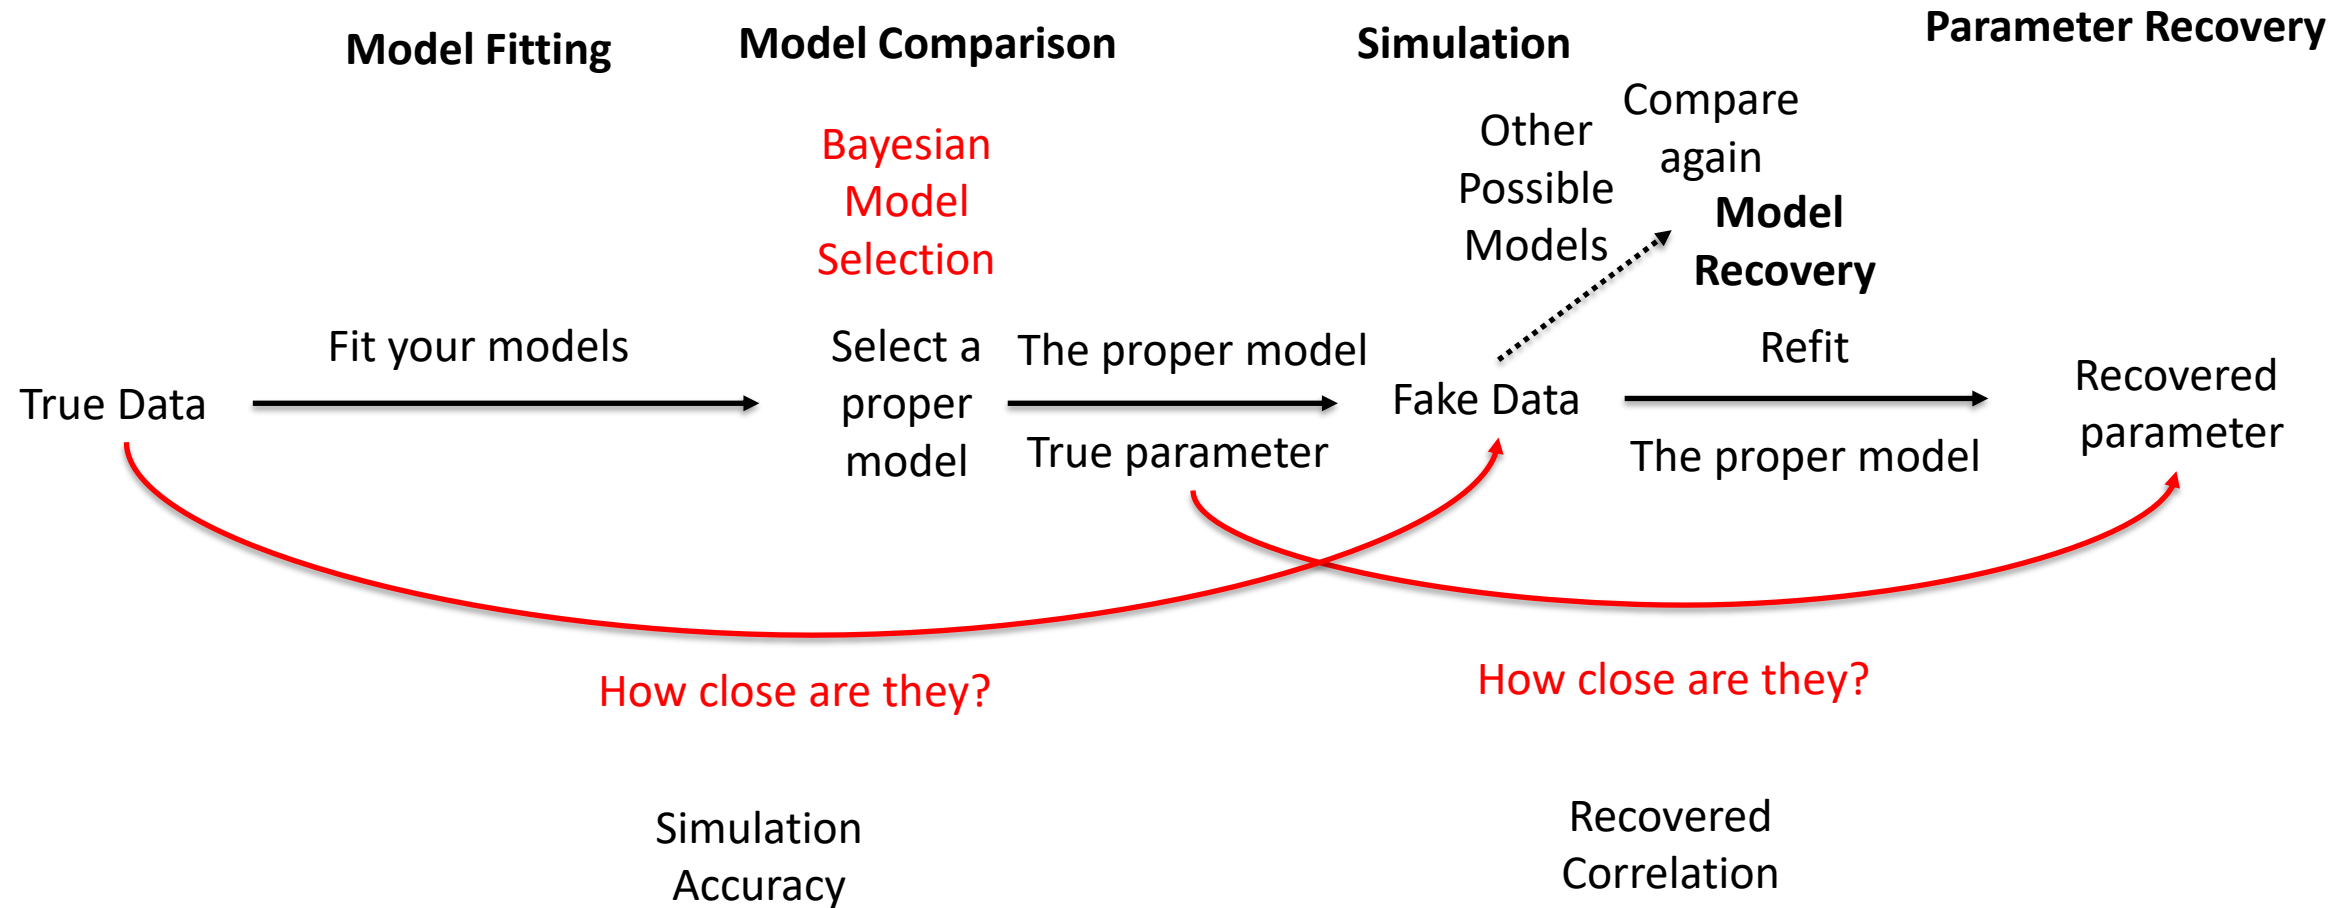




- If your model is fine with your experiment...
- Run the experiment!
- Why not run your model while collecting your data?



After you collect your data





- Your parameters:
 - Compare different parameters across your conditions;
 - Correlate your parameters with something you measure;
 - Even more flexible analysis approach (depends on your questions)
- Your latent variables:
 - Value
 - Prediction error/ TD error
 - Compare, correlate, cluster.....
 - Also more flexible analysis approach



- What your model (usually winning model) reflects about the cognitive, psychological or neuronal process?
- What your parameters mean? How are they bounded?
- Your fitting methods (specific details)
- Your parameter results (usually winning model); frequentist vs. Bayesian.
- Model Comparison (the criterion you choose, BMS result)
- How well is your model fitted, simulated(predicted) and recovered?
- How's your model relevant analysis conducted?
- How are the results above able to help answer your questions?



Three levels of description (David Marr, 1982)

Computational


Why do things work the way they do?
What is the goal of the computation?
What are the unifying principles?

Algorithmic

What representations can implement
such computations?
How does the choice of representations
determine the algorithm?



Implementational

How can such a system be built in
hardware?
How can neurons carry out the
computations?



maximize:

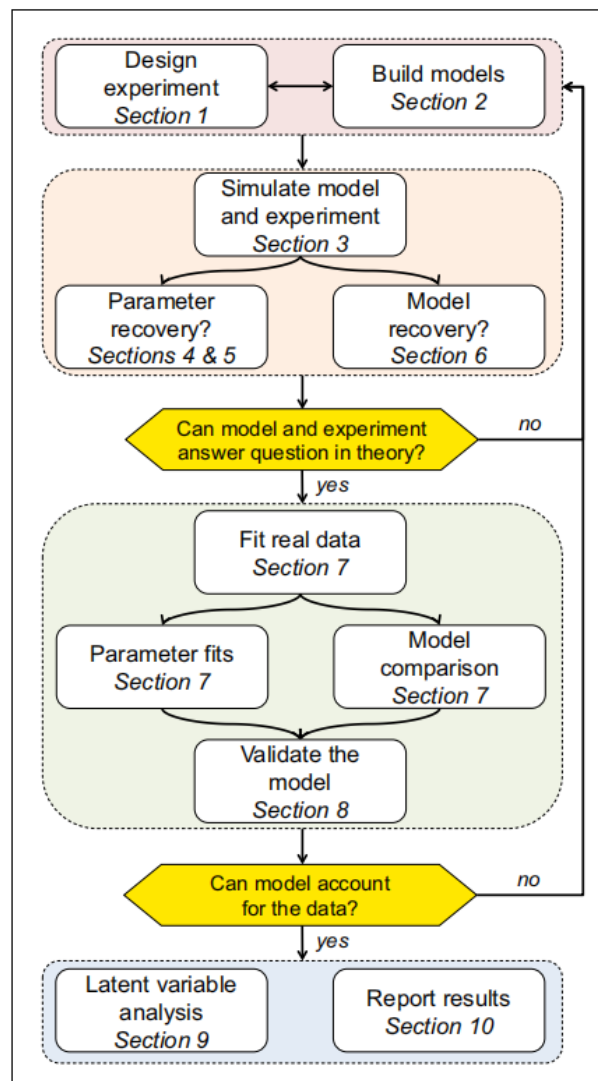
$$R_t = r_{t+1} + r_{t+2} + \dots + r_T$$



All models are wrong, **but**
some are useful.....

---- George E. P. Box

All models are **imperfect**,
but some are useful.....






MCMC and Stan



Monte Carlo Sampling



[A Zero-Math Introduction to Markov Chain Monte Carlo Methods](#)


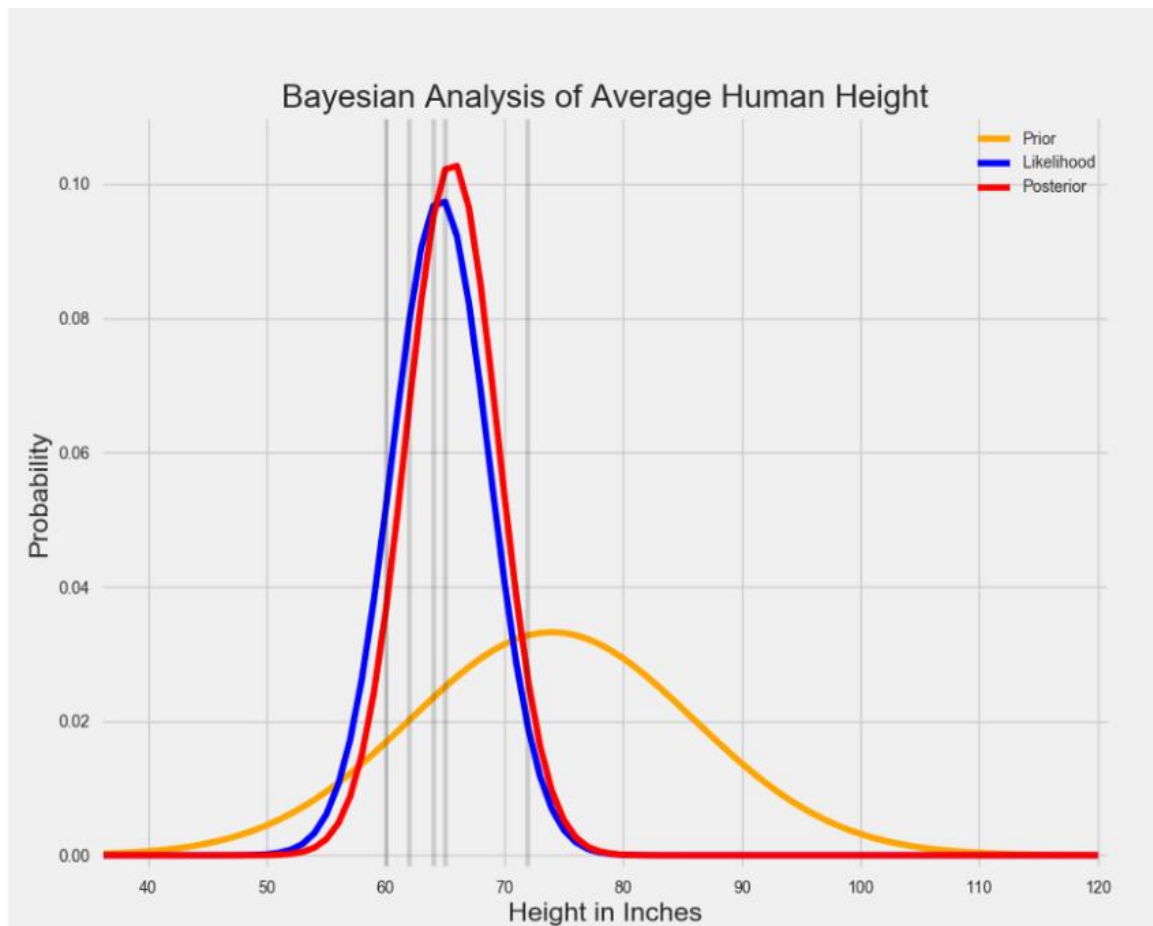


parameters data

$$P(\theta|D) \propto P(D|\theta)P(\theta)$$

Posterior Likelihood Prior

The diagram shows the equation $P(\theta|D) \propto P(D|\theta)P(\theta)$ with three red arrows pointing from the terms to their labels: from $P(\theta|D)$ to 'Posterior', from $P(D|\theta)$ to 'Likelihood', and from $P(\theta)$ to 'Prior'. Additionally, two red arrows point from the labels 'parameters' and 'data' to the terms θ and D respectively in the equation.






What is different from grid search?

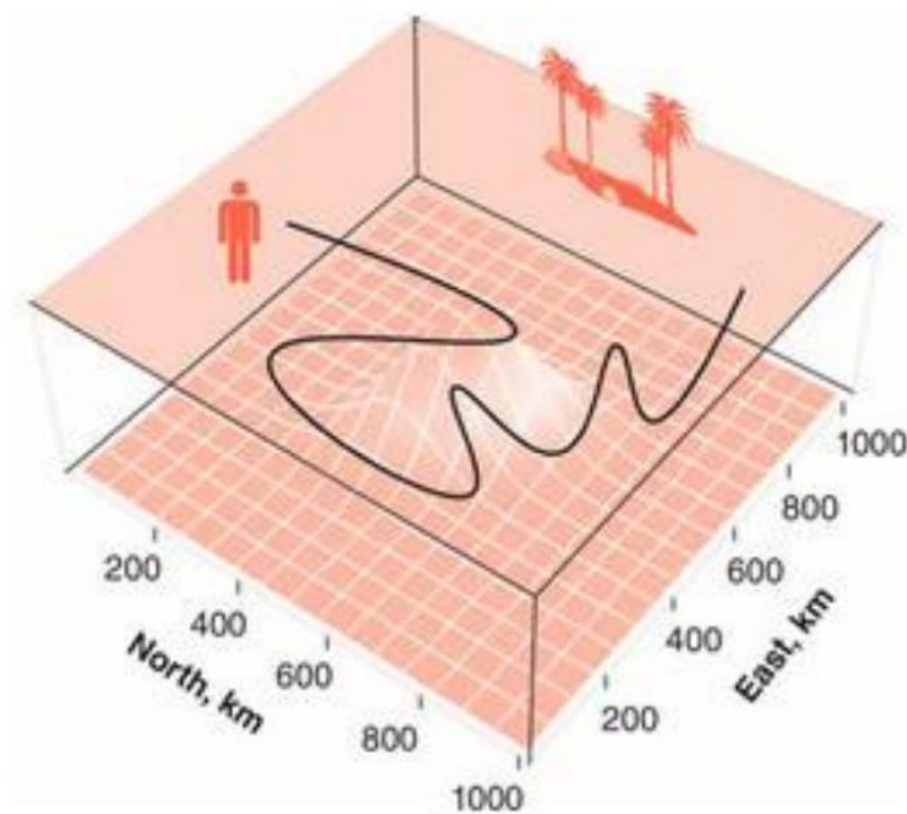


$$p(\theta | D) \propto p(D | \theta)p(\theta)$$





An MCMC Robot in 3D



Lambert (2018)



- Rejection Sampling
- Importance Sampling
- Metropolis Sampling
- Gibbs Sampling
- Hamilton MC (Stan)



- 1. an independent programming language based on C++
- 2. Can be interacted with different platforms (e.g., R, Matlab and Python)



- 1. Prepare your data
- 2. Write your model in Stan
- 3. Check with stan codes
- 4. Set up Stan options
- 5. Run/Call Stan files
- 6. Post-analysis



```
data {  
  //... read in external data...  
}
```

```
transformed data {  
  //... pre-processing of data ...  
}
```

```
parameters {  
  //... parameters to be sampled by HMC ...  
}
```

```
transformed parameters {  
  //... pre-processing of parameters ...  
}
```

```
model {  
  //... statistical/cognitive model ...  
}
```

```
generated quantities {  
  //... post-processing of the model ...  
}
```



Stan Syntax Cheatsheet.pdf

A minimal Stan program implementing a binomial model.

```
data {  
  int n;  
  int x;  
}  
parameters {  
  real<lower=0, upper=1> p;  
}  
model {  
  p ~ uniform(0, 1);  
  x ~ binomial(n, p);  
}
```



Running a Stan program is usually done from another language such as Python or R.

(Here assuming `model_string` contains the model from the last slide.)

```
library(rstan)
data_list <- list(n = 30, x = 10)
s <- stan(model_code = model_string,
          data = data_list)
```


```
import pystan
data_list = dict(n = 30, x = 10)
s = pystan.stan(model_code = model_string,
                data = data_list)
```




- 1. What is MCMC?
- 2. What is Stan?
- 3. How to use Stan?



Practices: implementing a reinforcement learning model in Rstan





- Agent perceives the environment state s
- Agent takes action a
- Agent receive reward r
- And agent adjust his action a'



$$PE = R - v_{i,t}$$

Prediction Error

$$v_{i,t+1} = v_{i,t} + \alpha * PE$$

α : learning rate

$$P(i, t + 1) = \frac{e^{v_{i,t+1} * \beta}}{\sum e^{v_{i,t+1} * \beta}}$$


Softmax function

$$\beta = 3^{\tau} - 1$$

Rescorla, R. A., & Wagner, A. R. (1972)



B



$P(\text{Reward})$

70%
30%

Uniform(8,12)



1

25

50
Trial

75

100

(Zhang, Glascher, 2020)




Task setting

```
##task setting (probability reversal learning task)
set.seed(2021)
reverse_trials=sample(c(8,9,10,11,12,13), replace = F)
ntrials=60
reward_prob=matrix(c(0.7,0.3,0.3,0.7),ncol=2,nrow=2)
```

Agent setting

```
##agent setting
nsubjects=60
lr=rnorm(n=nsubjects,mean=0.4,sd=0.2)
tau=rnorm(n=nsubjects,mean=3,sd=1)
```

```
#run the simulation
agent_data=list()
for (i in 1:nsubjects){
  tmp_lr=lr[i]
  tmp_tau=tau[i]
  v=matrix(rep(0,2*ntrials),nrow=2,ncol=ntrials)
  p=c(0,0)
  consistency=3**tmp_tau-1
  choice=rep(0,ntrials)
  reward=rep(0,ntrials)
  correct=rep(0,ntrials)
  n_reverse=1
  count_trial=0
  for (t in 1:ntrials){
    if(t>1){
      v[choice[t-1],t]<-v[choice[t-1],t-1]+tmp_lr*(reward[t-1]-v[choice[t-1],t-1])
    }
    p<-softmax(v[,t]*consistency)
    choice[t]=sample(c(1,2),prob = p,replace=T)
    reward[t]=sample(c(1,0),prob = reward_prob[choice[t],],replace=T)
    count_trial=count_trial+1
    if(count_trial==reverse_trials[n_reverse]){
      reward_prob=1-reward_prob
      count_trials=0
      n_reverse=n_reverse+1
    }
    if (reward_prob[choice[t],1]==0.7){
      correct[t]=1
    }else{
      correct[t]=0
    }
  }
  agent_data[[i]]=list(c(1:ntrials),choice,reward,correct,v,tmp_lr,tmp_tau)
}
```



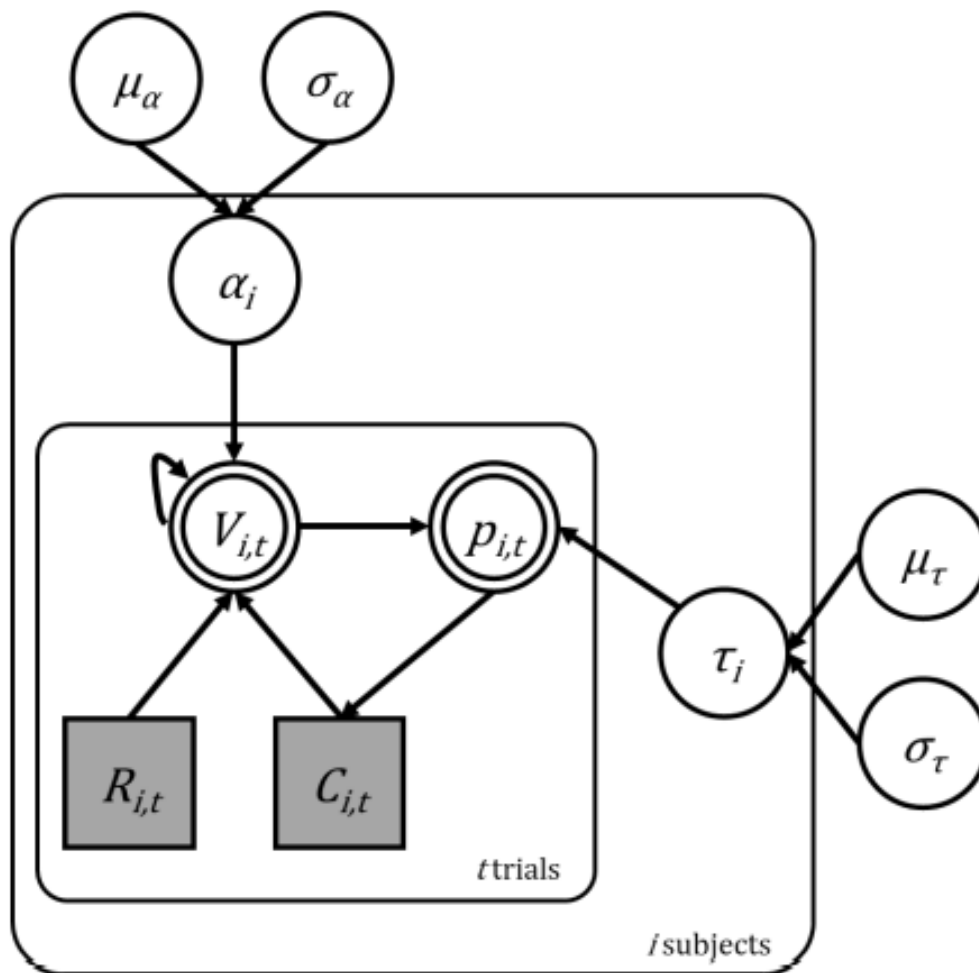




Define your data

```
data {  
  int<lower=1> nsubjects;  
  int<lower=1> ntrials;  
  int<lower=1,upper=2> choice[nsubjects, ntrials];  
  real<lower=0, upper=1> reward[nsubjects, ntrials];  
}
```

```
transformed data {  
  vector[2] initv; // initial values for v  
  initv = rep_vector(0.0, 2);  
}
```

Hierarchical Bayesian Model




Stan model

```
parameters {  
  // group-level parameters  
  real lr_mu_raw;  
  real tau_mu_raw;  
  real<lower=0> lr_sd_raw;  
  real<lower=0> tau_sd_raw;  
  
  // subject-level raw parameters  
  vector[nSubjects] lr_raw;  
  vector[nSubjects] tau_raw;  
}  
  
transformed parameters {  
  vector<lower=0,upper=1>[nSubjects] lr;  
  vector<lower=0,upper=5>[nSubjects] tau;  
  
  for (s in 1:nSubjects) {  
    lr[s] = Phi_approx( lr_mu_raw + lr_sd_raw * lr_raw[s] );  
    tau[s] = Phi_approx( tau_mu_raw + tau_sd_raw * tau_raw[s] ) * 5;  
  }  
}
```

Phi_approx: sigmoid

$$y = \frac{1}{1 + e^{-x}} \quad y \in (0,1)$$





```
model {  
  lr_mu_raw ~ normal(0,1);  
  tau_mu_raw ~ normal(0,1);  
  lr_sd_raw ~ cauchy(0,3);  
  tau_sd_raw ~ cauchy(0,3);  
  
  lr_raw ~ normal(0,1);  
  tau_raw ~ normal(0,1);  
  
  for (s in 1:nsubjects) {  
    vector[2] v;  
    real pe;  
    v = initv;  
  
    for (t in 1:nTrials) {  
      choice[s,t] ~ categorical_logit((pow(3,tau[s])-1)* v );  
      pe = reward[s,t] - v[choice[s,t]];  
      v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
    }  
  }  
}
```



```
generated quantities {
  real<lower=0,upper=1> lr_mu;
  real<lower=0,upper=3> tau_mu;

  real log_lik[nSubjects];

  int y_pred[nSubjects, nTrials];

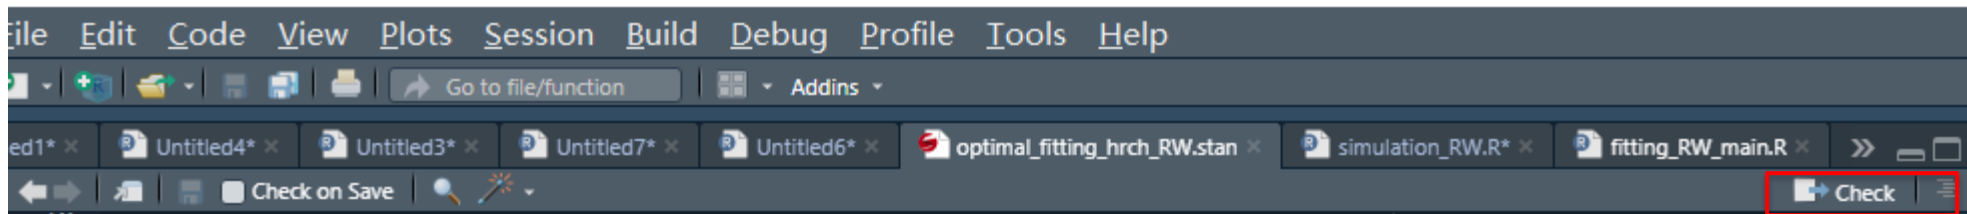
  lr_mu = Phi_approx(lr_mu_raw);
  tau_mu = Phi_approx(tau_mu_raw) * 5;
  y_pred = rep_array(-999,nSubjects ,nTrials);

  { // local block
    for (s in 1:nSubjects) {
      vector[2] v;
      real pe;
      log_lik[s] = 0;
      v = initv;

      for (t in 1:nTrials) {
        log_lik[s] = log_lik[s] + categorical_logit_lpmf(choice[s,t] | ((pow(3,tau[s])-1) * v));
        y_pred[s,t] = categorical_logit_rng( (pow(3,tau[s])-1) * v );
        pe = reward[s,t] - v[choice[s,t]];
        v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;
      }
    }
  }
}
```



- Check your codes



```
> rstan::rstudio_stanc("D:/Hanbo/psychology/科研/实验室申请/北大周晓林实验室/modeling 教学/codes/optimal_fitting_hrch_RW.stan")
D:/Hanbo/psychology/科研/实验室申请/北大周晓林实验室/modeling 教学/codes/optimal_fitting_hrch_RW.stan is syntactically correct.
> |
```

- Bugs?

[The Stan Forums \(mc-stan.org\)](http://mc-stan.org)



- Prepare data
- Specify stan file
- Set options for fitting
- Run the fitting process



Original data: list \longrightarrow Target data: array

```
#suppose we have completed simulation process and get the agent_data

nSubjects <- length(agent_data)
nTrials   <- length(agent_data[[1]][[1]])

##reshape the data from a list to an array to be fitted
data_array<-array(rep(0,nSubjects*nTrials*2),dim=c(nSubjects, nTrials, 2))
for (i in 1:nSubjects){
  for(t in 1:nTrials){
    data_array[i,t,1]=agent_data[[i]][[2]][t]
    data_array[i,t,2]=agent_data[[i]][[3]][t]
  }
}

dataList <- list(nSubjects=nSubjects,
                nTrials=nTrials,
                choice=data_array[, ,1],
                reward=data_array[, ,2])
```



Your
own
path

```
#####
rstan_options(auto_write = TRUE)
options(mc.cores = 4)

modelFile <- 'D:\\Hanbo\\psychology\\科研\\实验室申请\\北大周晓林实验室\\modeling 教学\\codes\\optimal_fitt

nIter      <- 16000
nChains     <- 4
nWarmup    <- floor(nIter/2)
nThin      <- 1

cat("Estimating", modelFile, "model... \n")
startTime = Sys.time(); print(startTime)
cat("calling", nChains, "simulations in stan... \n")

fit_rl <- stan(modelFile,
               data      = dataList,
               chains    = nChains,
               iter      = nIter,
               warmup    = nWarmup,
               thin      = nThin,
               init      = "random",
               control    = list(adapt_delta=0.999, max_treedepth=100),
               seed      = 2021
               )

cat("Finishing", modelFile, "model simulation ... \n")
endTime = Sys.time(); print(endTime)
cat("It took", as.character.Date(endTime - startTime), "\n")
```




```
run_RW_fitting <- function() {  
  # =====  
  ##### Construct Data #####  
  # =====
```

Run the whole function to define the function

```
+ ]  
> fit_rl<-run_RW_fitting()
```

Call the function through command



```
Environment History Connections Tutorial
Files Plots Packages Help Viewer

Click the Refresh button to see progress of the chains
starting worker pid=15180 on localhost:11696 at 15:45:07.415
starting worker pid=2264 on localhost:11696 at 15:45:07.572
starting worker pid=6536 on localhost:11696 at 15:45:07.742
starting worker pid=11176 on localhost:11696 at 15:45:07.899

SAMPLING FOR MODEL 'optimal_fitting_hrch_RW' NOW (CHAIN 1).

SAMPLING FOR MODEL 'optimal_fitting_hrch_RW' NOW (CHAIN 2).

SAMPLING FOR MODEL 'optimal_fitting_hrch_RW' NOW (CHAIN 3).
Chain 1:
Chain 1: Gradient evaluation took 0.004 seconds
Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 40 seconds.
Chain 1: Adjust your expectations accordingly!
Chain 1:
Chain 1:
Chain 1: Iteration:      1 / 16000 [ 0%] (Warmup)

SAMPLING FOR MODEL 'optimal_fitting_hrch_RW' NOW (CHAIN 4).
Chain 2:
Chain 2: Gradient evaluation took 0.003 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 30 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration:      1 / 16000 [ 0%] (Warmup)
Chain 3:
Chain 3: Gradient evaluation took 0.003 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 30 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:      1 / 16000 [ 0%] (Warmup)
Chain 4:
Chain 4: Gradient evaluation took 0.004 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 40 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration:      1 / 16000 [ 0%] (Warmup)
```




```
> print(fit_r1)
Inference for Stan model: optimal_fitting_hrch_RW.
4 chains, each with iter=16000; warmup=8000; thin=1;
post-warmup draws per chain=8000, total post-warmup draws=32000.
```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
lr_mu_raw	-0.50	0.00	0.18	-0.87	-0.62	-0.49	-0.38	-0.16	2701	1
tau_mu_raw	1.83	0.00	0.42	1.12	1.53	1.79	2.08	2.77	11334	1
lr_sd_raw	1.06	0.00	0.16	0.78	0.94	1.05	1.16	1.42	3042	1
tau_sd_raw	2.00	0.01	0.61	1.03	1.58	1.93	2.34	3.41	4095	1
lr_raw[1]	-0.91	0.01	1.18	-2.78	-1.82	-1.08	-0.07	1.55	9682	1
lr_raw[2]	0.69	0.00	0.36	0.05	0.44	0.67	0.92	1.46	14462	1
lr_raw[3]	-0.22	0.00	0.46	-1.15	-0.54	-0.21	0.10	0.65	17157	1
lr_raw[4]	0.49	0.00	0.35	-0.13	0.26	0.47	0.71	1.24	13418	1
lr_raw[5]	1.84	0.00	0.49	0.85	1.52	1.85	2.16	2.78	13315	1
lr_raw[6]	-0.27	0.00	0.50	-1.28	-0.60	-0.25	0.09	0.64	19402	1
lr_raw[7]	0.49	0.00	0.34	-0.10	0.26	0.46	0.70	1.25	12884	1
lr_raw[8]	-0.70	0.00	0.28	-1.25	-0.88	-0.69	-0.51	-0.15	3689	1
lr_raw[9]	0.63	0.00	0.45	-0.42	0.39	0.69	0.93	1.39	12864	1
lr_raw[10]	1.57	0.01	0.68	0.24	1.08	1.63	2.06	2.81	12984	1
lr_raw[11]	-1.31	0.01	0.37	-2.05	-1.56	-1.31	-1.07	-0.61	3577	1
lr_raw[12]	0.77	0.01	0.75	-0.43	0.23	0.63	1.29	2.35	7887	1
lr_raw[13]	0.61	0.00	0.43	-0.13	0.31	0.57	0.86	1.57	15384	1
lr_raw[14]	0.36	0.01	0.70	-1.16	-0.10	0.48	0.87	1.46	18117	1
lr_raw[15]	0.53	0.00	0.28	0.01	0.33	0.51	0.70	1.13	9824	1

Rhat should generally be smaller than 1.01

[‘shinystan’](#)





Stan Sampling Parameters

cognitive model

statistics


computing

parameter	description	constraint	default
iterations	number of MCMC samples (per chain)	int, > 0	2000
delta: δ	target Metropolis acceptance rate	$\delta \in [0, 1]$	0.80
stepsize: ε	initial HMC step size	real, $\varepsilon > 0$	2.0
max_treedepth: L	maximum HMC steps per iteration	int, $L > 0$	10

Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max_treedepth

```
funnel_fit2 <- stan("_scripts/funnel.stan",  
  iter = 4000,  
  control = list(adapt_delta = 0.999,  
    stepsize = 1.0,  
    max_treedepth = 20))
```





- Mean
- HDI=highest density interval (credential interval)

```
#' Compute Highest-Density Interval
#
#' @description
#' Computes the highest density interval from a sample of representative values, estimated as shortest credible interval
#' Downloaded from John Kruschke's website \url{http://www.indiana.edu/~kruschke/DoingBayesianDataAnalysis/}
#
#' @param samplevec A vector of representative values from a probability distribution (e.g., MCMC samples).
#' @param credMass A scalar between 0 and 1, indicating the mass within the credible interval that is to be estimated
#
#' @return A vector containing the limits of the HDI
#
#' @export

HDIofMCMC = function(samplevec,
                      credMass = 0.95 ) {
  if ( class(samplevec) == "mcmc.list" ) {
    samplevec = as.matrix(samplevec)
  }
  sortedPts = sort( samplevec )
  ciIdxInc = floor( credMass * length( sortedPts ) )
  nCIs = length( sortedPts ) - ciIdxInc
  ciwidth = rep( 0 , nCIs )
  for ( i in 1:nCIs ) {
    ciwidth[ i ] = sortedPts[ i + ciIdxInc ] - sortedPts[ i ]
  }
  HDImin = sortedPts[ which.min( ciwidth ) ]
  HDImax = sortedPts[ which.min( ciwidth ) + ciIdxInc ]
  HDIlim = c( HDImin , HDImax )
  return( HDIlim )
}
```



北京大学心理与认知科学学院

School of Psychological and Cognitive Sciences, Peking University

How to analyze?

Statistical Rethinking

A BAYESIAN COURSE
WITH EXAMPLES
IN R AND STAN

Richard McElreath



Generally, you can compare your models by using various indicators:

- Log-likelihood
- Alike Information Criterion (AIC)
- Bayesian Information Criterion (BIC)
- **Leave-one-out Information Criterion (LOOIC)**

```
##model comparison  
library(loo)  
loo(fit_r1)
```

```
Computed from 32000 by 60 log-likelihood matrix
```

	Estimate	SE
elpd_loo	-879.4	104.2
p_loo	62.8	4.2
looic	1758.8	208.5

```
-----  
Monte Carlo SE of elpd_loo is NA.
```



- 1. How to write a stan code?
- 2. How to write an R code?
- 3. How to diagnose the fitting result?
- 4. How to analyze the result?



- We introduce concepts of computational modeling and why we need to learn about it in psychology and cognitive sciences.
- We introduce standard approaches (ten simple rules) on conducting a computational modeling research.
- We briefly go through how Stan works with MCMC.
- We dived into the codes to realize from the very beginning of simulation to analyze model fitting results.



- How to create/think up a good model?
 - Sorry, I don't know either...
- How to learn about computational modeling?
 - Knowledge perspective:
 - Math (linear algebra, theory of probability, Calculus)
 - Programming languages (R, Matlab, Python, Julia, C++, Java.....)
 - Algorithms (basics about machine learning)
 - Practice perspective:
 - Learn by demanding, doing and achieving!



- Summer schools:

- [Neuromatch Academy Summer School \(Annually\)](#)
- [Computational Psychiatry Course @ Zürich, CH \(annual\)](#)
- [London Computational Psychiatry Course @ London, UK \(annual\)](#)
- [Brains, Minds & Machines Summer Course @ MIT, US \(annual\)](#)
- And so on (keep searching!)



- Online courses:
 - [Machine learning \(@Coursera- Andrew Ng\)](#)
 - [Bayesian Modeling \(@Bilibili- Lei Zhang\)](#)
 - [Reinforcement learning model basics \(myself\)](#)
 - [Reinforcement learning \(@coursera\)](#)



- Books:
 - R语言实战
 - Statistical Rethinking
 - 认知和行为的计算建模



Acknowledgement



Thank Dr. Zhou for his supervision and my lab members for their support!

Special thanks to:

- Wenqian Lin (RA at Beijing Normal University)
- Shen Xu (PhD student at Peking University)

For their advices on the tutorial.

Thank you for your listening!